

CareStitch

A well-built scheduling workspace from people who have done the job: capacity in points by day on a map, territories drawn to the street, a staffing-request workflow that replaces email and Teams blasts, and care-team chat attached to every patient. It shows schedulers everything and decides nothing. There is no HCHB integration in production and no patient-facing automation.

SCORECARD GOING IN 56 Conditional · hold · no open stop-checks	LIVE FOOTPRINT 79 orgs 69% home health, 144 branches. Largest 5.3k, in one branch	GUIDE QUESTIONS 4 / 18 answered in full; 8 partly; 2 not answered; 4 not asked	NEXT TOUCHPOINT Follow-up Virtual call on HCHB integration and automation R&D, before the on-site
---	--	---	--

Six things that matter

- 01 Built by schedulers, for schedulers.** The COO is a home health PT and the operations manager ran centralized HCHB scheduling for 10 years. Priority groups for requests, a call log, and drive-time blocks on the clinician's calendar are details only field people would build. **0:00:53**
- 02 Capacity is visible, not calculated for you.** Points by clinician by day, colour-coded, on a map. A branch leader can enter an address and discipline and see who could take a new referral. There is no branch-level total and no forecast. **0:43:13**
- 03 The request workflow is the product's best feature.** Priority groups (full-time first, per diem after a delay), reminders, a timestamped read-and-response log, clinician comments, and an optional auto-assign to the first taker. It replaces the Teams-and-email scramble with data. **1:16:23**
- 04 It advises. People still do the clicking.** Evan named it on the call: a great dashboard, but every action is manual. CareStitch agreed. Automation and agentic AI are "research mode," and a scheduling recommendation engine must come first. **1:56:16**
- 05 HCHB integration is not in production.** The questionnaire return describes a read-only beta by screen automation, with no write-back. The demo didn't cover it, and Evan listed it as the first topic for a follow-up. Home health guardrails (SOC-before-eval, LUPA) "live in the EMR." **1:44:09**
- 06 Patient engagement is roadmap.** Clinicians still phone patients, now with click-to-call and an automatic call log. Automated confirmation is a top-five roadmap item; rescheduling and negotiation wait on the recommendation engine. **1:10:06**

SUGGESTED SCORECARD MOVEMENT


Hold at Conditional. A strong workspace, not yet the engine we are buying.

The demo confirmed the return: excellent visibility and coordination for schedulers, built on real home health experience. Integration, automation and patient engagement are all still ahead. At Compassus scale the value case rests on automation, which CareStitch has not started building.

Take the follow-up call they offered. Advance only if it shows a dated HCHB production plan (read and write), a working prototype of recommendations, and the automation R&D they offered to show.

The demo showed that the dashboard and request workflow patterns are worth copying, whoever we pick.

The shape of the call




CareStitch opened with discovery about our intake build, branch variation and priorities, then ran a scenario-based demo of about 80 minutes. The integration discussion never happened and was deferred to a follow-up.

Our position, as stated in the room

ITEM	AS STATED	WHEN
Scale	About 10,000 nurses across the enterprise; about 340 schedulers	0:03:25 · 1:52:21
Intake platform	Custom build with an engineering firm (Commure, per the Axle call). Outputs a complete, valid referral; written to HCHB by RPA; open to flat files and APIs both ways. Capacity at accept/decline deliberately left out of scope for this project	0:08:25 · 0:13:49
Capacity today	Grids and a SharePoint site updated by hand; spreadsheets and sticky notes vary by branch	0:13:05 · 0:16:51
Scheduling model	Local today; open to centralized or hybrid. Standardization across branches is part of the business case	0:14:33 · 0:17:34
Documentation	Automated; clinicians barely work in PointCare now	1:03:32
Priorities	An industry-specific thought partner (Evan); the clinician experience (Colin); adaptability and co-development (Michael)	0:18:54
Biggest ask	Take the 30-60 min a day of patient calls off 10,000 clinicians, with agentic AI if possible	1:07:57

Configure, see capacity, assign or request, coordinate

CareStitch separates the *task*, the work the EMR requires, from the *visit*, the appointment that fulfils it. The office assigns tasks, directly or through a request, and the clinician books the visit and phones the patient. HCHB stays the record for documentation and completion.



● LIVE demonstrated in software

● HUMAN-LED supported, but a person does the work

● NOT BUILT named as roadmap or absent

Everything the agency knows, made into filters

WHAT IT DOES


- **Provider types** (RN, LPN, PT, PTA, OT, COTA, MSW, HHA, SLP and others) and **visit types** with point values, deadline and threshold hours. For example, SOC is 2 points with a 48-hour deadline and ROC is 1.5.
- **General visits** such as meetings, case conferences and breaks block clinician time, with a Teams link attached if wanted.
- **Custom properties**, self-service: patient labels shown to the care team (e.g. “hospitalized” appears on every care-team phone at once), or matching properties set on both patient and clinician (PICC line, facility, language) that filter out ineligible clinicians automatically.
- Facility matching lets an agency keep two nurses per building for continuity and the facility relationship.

FOR COMPASSUS


- Michael asked about labels set automatically from data flowing in. That is live for some EMRs’ patient-flag columns, and their API (in development) could accept labels from our intake platform. 0:28:28
- Colin saw emergency-planning value: find at-risk patients by area in hurricanes or fires. Their clients have done this for wildfires.
- All of it is configured by agency administrators with no vendor involvement. Governance across 80+ branches is on us.

Code	Name	Deadline Hrs:AM	Threshold Hrs:AM	Color	Points	Auto-fill Address	Billable
SOC	Start of Care	48:00	24:00	Green	2	☒	☒
ROC	Resumption of Care	48:00	24:00	Teal	1.5	☒	☒
EVAL	Evaluation	96:00	24:00	Orange	1.5	☒	☒
FUV	Follow Up Visit	96:00	24:00	Blue	1	☒	☒
SUP	Supervisory Visit	48:00	24:00	Dark Blue	1	☒	☒
RAS	Reassessment	96:00	24:00	Dark Blue	1.5	☒	☒
RCT	Recertification	96:00	24:00	Teal	1.5	☒	☒
DC	Discharge	96:00	24:00	Orange	1	☒	☒
ODC	Oasis Discharge	96:00	24:00	Pink	1.5	☒	☒
TELE	Telehealth	48:00	24:00	Blue	1	☐	☒

0:24:46 Visit types: points, deadline and threshold hours, auto-fill address, billable



0:24:10 Provider types, with colour and category



0:27:43 Custom properties: patient labels and matching properties



0:32:16 “Facilities” used to keep the same nurses in each building


Territories drawn to the street, not the ZIP

WHAT IT DOES


- Regions are built from ZIP codes or **drawn polygons**, to split a big ZIP, follow a highway, or carve up a dense city.
- Regions can overlap; a clinician can belong to several; or a region can be one clinician's personal territory (useful for choosy per diems).
- Once set, the system knows a patient's region from the address and which clinicians serve it. Nobody needs to carry territory knowledge in their head.

FOR COMPASSUS

- Michael raised **Napa**, where one town, ZIP and county forces Google-mapping every patient. Polygons fix that.
- Colin asked for census and caseload overlays on territories. They were promised on the map view but not shown as a territory-level analysis.
- Built-in territory knowledge is what makes a centralized scheduling model workable.



0:36:25 Region management: overlapping territories across Nashville and Williamson County



0:39:13 Regions, with the clinicians who serve them


Can we take this referral? Look at the grid

WHAT IT DOES


- The schedule calendar shows each clinician by row and the Medicare week Sunday to Saturday, with **points still open** in bold. Grey marks days off. The map shows clinician homes and this week's patients.
- Each clinician has expected points per day. Assigning a visit deducts its points, so "no more mental gymnastics."
- To check a referral: enter the address, filter by discipline, employment type and properties (e.g. PICC line), and see who has room on which day.
- Shapes on the grid mark whether a task has been booked as a visit.

FOR COMPASSUS


- This answers common question 1 at the clinician level. There is **no branch total, no forecast and no referral-inflow view**. The leader adds it up by eye.
- Michael wants to query this from the intake tool: "do we have someone who can see this patient?" CareStitch's API is in development. 0:47:11
- Pending-auth demand can only be shown with a custom property; auth isn't ingested.




0:43:31 Referral check: enter an address, filter to full-time and part-time RNs with a PICC competency, see open points by day




0:24:01 Schedule mode: clinicians by day with open points, and the map



0:42:46 A clinician's work schedule: expected points per day, days off



0:45:04 Filters: provider type, employment, region, language, properties



0:45:10 Distance filters and minimum available points


Direct assignment with the context in view

WHAT IT DOES


- Select the patient and the disciplines needed. The grid shows each eligible clinician's distance for each day, from this patient to their nearest visit that day or to their home.
- A yellow asterisk marks a clinician who has seen the patient before, for continuity.
- The scheduler picks, the tasks go to the clinician, and a push notification reaches their phone.

FOR COMPASSUS

- **No recommendation or ranking.** The system filters and displays; the scheduler judges. Their return calls this "advise and orchestrate."
- Evan asked whether plan-of-care data from intake could fill in the disciplines and visits automatically. It would take the future API, and today it is clicks. 0:48:06



0:47:13 New patient: RN and PT candidates, with distance and open points by day



0:49:25 Assigning the SOC to the nurse who has seen the patient before


The staffing request: an offer that finds its taker

WHAT IT DOES


- Build a request: patient, visit type, deadline (e.g. SOC within 48 h), discipline, gender or language preference, notes (“dialysis, afternoons only”).
- Eligible clinicians are sorted into **priority groups** (full-time, part-time, per diem, contractor), each with a configurable delay and reminder. The scheduler can deselect someone (“in documentation jail”) or move them to a later group.
- Clinicians see the visit, its distance from their home and from each booked visit, and the deadline. They reply available or not available and can add a comment (“only 2pm Friday”).
- A queue colours each request: red when the deadline is close and nobody has said yes, blue when someone is available. A timestamped log shows who was sent it, who read it and who answered.
- Assignment can be manual (favour full-time or closest) or automatic to the first clinician available. If nobody takes it, add recipients from other territories.

FOR COMPASSUS


- This replaces the Teams, email and text scramble with **data we don't have today**: response speed and accept rates by clinician.
- Colin: branches cherry-pick visits in Teams that should go together. This structure stops that.
- Colin saw uses beyond visits: on-call, therapy weekend rotations. Their clients already do this with staffing requests.
- The first auto-assign version was reworked because agencies wanted control. Pick the default deliberately.




1:19:07 User selection: priority groups with send delays and energy bars




1:16:55 Step 1: request type, deadline, clinician preference, notes



1:16:34 Request Management: red is at risk, blue has a taker



1:26:46 Request details and a timestamped log of every read and reply



1:31:10 The clinician's request card: distance, deadline, available or not


The clinician still calls, but with better tools

WHAT IT DOES


- The mobile app home shows today's appointments and a productivity summary (points to target). Patients are sorted by the next task due.
- The **patient-centred calendar** shows the clinician's other appointments with distance from this patient, a **drive-time** overlay that greys out travel time, and other disciplines' visits for this patient in red, so the PT sees the SOC is Thursday.
- Click-to-call any number on file, with an **automatic call log** and notes the office can see.
- The office side has an Optimize Route button for centralized scheduling. Clinicians place visits one at a time around ones already confirmed, anchoring the fixed ones first.

FOR COMPASSUS


- Automated confirmation is a top-five roadmap item**, and text comes before voice. Rescheduling and negotiation need the recommendation engine first. No dates. **1:10:06**
- Colin asked about calls showing the branch number, since patients ignore unfamiliar area codes. It goes through RingCentral-type integrations; flagged for follow-up.
- Evan: documentation is automated, so clinicians are barely in PointCare. That could make CareStitch their single screen.




1:04:37 Patient-centred calendar: drive time greyed, other disciplines in red




0:23:16 Office portal and clinician app side by side




0:52:37 Clinician calendar with distance to each booked visit



0:55:58 Map view: booked visits and patients not yet booked



0:58:40 Booking the visit, with a reminder note ("call DPOA first")



1:10:07 Call log, visible to the office and marketing


A call-out, reassigned by hand

WHAT IT DOES


- Open the clinician who called out and reassign each visit. Eligible clinicians are pre-filtered by discipline and enough points for the visit.
- Widen the search to PTAs; see each candidate's distance to the patient and to their booked visits; save.
- If nobody fits, send it as a request (stage 04). Clinicians can also start a reassignment request from their phone.

FOR COMPASSUS

- It works and it is quick, but it is manual per visit, with no suggested replacement, no ranking and no bulk accept.
- **No home health guardrails.** Asked about SOC-before-eval when dates shift: "we don't currently have those guardrails... the guardrails live in the system of record." LUPA is roadmap. 1:44:09



1:15:34 Reassign: PTAs added, each with distance to the patient and their day's visits



1:14:31 The clinician who called out, with visits to move


Every patient gets a care-team room

WHAT IT DOES


- A secure chat room is created for each new patient, and every clinician assigned joins it automatically. The office is included, and newcomers see the full history.
- Channels: broadcast (office speaks, field reads) and restricted (clinicians post, only the office sees). Plus direct messages and groups.
- Chat history exports per certification period, for interdisciplinary-communication evidence in audits.
- **Sub-organizations:** a clinician working for several branches or lines shows one combined availability.
- Staffing requests double as a ticket queue for pending admissions and other office work.

FOR COMPASSUS


- This is the strongest coordination layer of the three vendors. Compare with Teams, where the scramble happens today.
- Evan asked for scheduler productivity reporting (who touched what, how long). Not built; “never been asked.” 1:52:21




1:39:49 A patient's care-team chat, with a wound photo



1:39:16 Patient thread with care-team messages



1:42:16 Broadcast channels on the clinician's phone



1:38:28 Staffing requests as a queue for pending admissions

Impact and positioning

CLAIM	AS STATED	OUR READ
Scheduler capacity	● 2× CENSUS One customer doubled census without adding a scheduler	Self-reported, no baseline. Encouraging for our centralization case.
Mileage	● -1.27 MI/VISIT Over 104k visits, about \$92.7k a year; “\$1 in, \$3 back” from mileage alone	The only quantified result, from one customer and not validated. Far below Axle’s claimed –32%.
Adoption	Median 93% monthly clinician activation by month six (return)	Strong, and measured across the whole roster.
Uptime	99.999% of requests served, no outage in 12 months; offline-first mobile; 99.9% contractual	The best continuity answer in the field.
Data maturity	“Almost none of them really are that data driven”	Candid. It also means few provable outcomes to point to.
Partnership	Design partner, joint governance, phased and volume pricing, development credits (return)	The most structured partnership offer so far. Not discussed on the call.

Question by question

● ANSWERED 4 ● PARTLY 8 ● NOT ANSWERED 2 ● NOT ASKED 4

QUESTION	STATUS	WHAT WE HEARD	WHEN
Common 01 The referral decision	● ANSWERED	Shown live: address, discipline and property filters give open points by clinician and day. Clinician level only, with no branch total or forecast.	0:43:13
Common 02 The episode, not the visit	● NOT ANSWERED	A Medicare-week grid and visit deadlines, but no order ingestion, pacing, LUPA or recert logic. “Guardrails live in the system of record.”	1:44:09
Common 03 HCHB freshness, with a number	● NOT ASKED	Deferred to a follow-up. Their return: read-only beta by screen automation, configurable cadence, not in production.	1:55:31
Common 04 Advise versus decide	● ANSWERED	Advise and orchestrate; the scheduler decides. Auto-assign to the first taker was the first version and was rolled back to agency choice. Agentic AI is research.	1:29:13
Common 05 The clinician says no	● PARTLY	Mark not available, with a comment; request a reassignment from the phone. Responses are logged as data. No model learns from it.	1:31:11
Common 06 A Tuesday that goes wrong	● ANSWERED	Live: reassign a call-out’s visits by hand; if nobody fits, a request with priority groups; if nobody takes it, widen to other territories.	1:14:24
Common 07 What you don’t do	● PARTLY	Named: patient confirmation, recommendation engine, API, EMR integration (priorities), LUPA and sequence guardrails. No dates.	1:10:06
Common 08 What breaks at our size	● PARTLY	Largest is 5.3k in one branch; 2.2k across six. Nothing said about what broke.	1:48:48
Common 09 The one number you’d stake it on	● PARTLY	Doubled census without a new scheduler; mileage \$1 in, \$3 back. Anecdotal: “we wish we had more information.”	1:45:12
Common 10 Where you’ll look worse	● NOT ASKED	Evan’s closing point, that it is all manual, drew the honest answer: automation is the next stage and still in research.	1:56:16

QUESTION	STATUS	WHAT WE HEARD	WHEN
CareStitch 01 HCHB dates for read and write	● NOT ASKED	Not covered. Named as the first follow-up topic.	1:55:31
CareStitch 02 The manual-upload path; HCHB customers	● PARTLY	Natalie Bell ran HCHB scheduling for 10 years. Tasks come “from the EMR”; completion syncs back as a check mark. No HCHB customer count given.	1:00:41
CareStitch 03 Orders, auth, compliance windows	● NOT ANSWERED	Not ingested; visit deadlines exist, and auth can only be a custom property. No enforcement.	1:44:09
CareStitch 04 “Advise and orchestrate”: principle or stage?	● ANSWERED	A stage. Stage one was “make the data available”; next is automation, then a recommendation engine and agentic AI, with the agency keeping approval.	1:56:16
CareStitch 05 Three named home health references	● NOT ASKED	Not requested on the call.	—
CareStitch 06 Will the product pace the episode?	● PARTLY	LUPA is roadmap; pacing stays with the scheduler for now.	1:44:09
CareStitch 07 Show the 69% home health concentration	● PARTLY	Domain depth was evident in the team and workflows. Customer names and EHRs not given.	0:21:34
CareStitch 08 Company size and funding	● PARTLY	Growth by word of mouth and a few conferences a year. Size and funding not given.	1:45:53

What to weigh

HIGH

Automation is not built

At 340 schedulers and 10,000 clinicians the value is in automation. CareStitch has the workspace and is only researching what goes on top of it.

HIGH

HCHB integration

A read-only beta by screen automation that has never run in production, with no write-back. It was not discussed on the call.

HIGH

No home health logic

No order, auth or frequency ingestion; no SOC-before-eval, LUPA or window enforcement. The scheduler and the EMR carry it all.

MEDIUM

Patient calls stay with clinicians

Our biggest single ask is roadmap: confirmation first, negotiation later, no dates. Click-to-call and a call log help, but they don't remove the work.

MEDIUM

Clicks at scale

Every assignment and reassignment is manual. What helps one scheduler becomes a cost across 340.

MEDIUM

Scale and proof

The largest customer has 5.3k patients in one branch; most outcomes are anecdotal. We would be the enterprise proof point.

MEDIUM

Configuration sprawl

Self-service properties and priority groups are powerful, and across 80+ branches they will drift without central governance. That fits our standardization goal only if we enforce it.

IN THEIR FAVOUR

Uptime and offline mobile

No outage in 12 months; clinicians keep the day's work when offline. The best continuity story of the three.

IN THEIR FAVOUR

Field-built, coordination-first

The request workflow, care-team chat, drive-time and call log are the best coordination tooling seen so far. The team knows HCHB and the branch reality.

What happens now

Agreed on the call

1. BOTH

A follow-up virtual call before any on-site: HCHB integration requirements and data refresh cadence, and CareStitch's automation R&D.

2. CARESTITCH

Look into scheduler productivity reporting (who touched what, how long).

3. COMPASSUS

Colin's note on the on-site schedule change shared with Shelley; she will pass it to her team.

4. COMPASSUS

Consider calls showing the branch number (RingCentral-type integration) as a separate topic.

Recommended for the follow-up

- HCHB: a dated plan for read and write, and the method (API, HL7 or RPA); what the beta reads; target cadence; HCHB customer count.
- The automation R&D they offered to show, and a scope and date for the recommendation engine.
- Patient confirmation by text: scope, date and a design partnership on the agentic version.
- Run the shared scenario through a 60-day episode: frequency, LUPA, recert, and SOC-before-eval when dates move.
- An API to query capacity from our intake platform and push patients and labels in.
- Three named home health references, a 2k+ ADC multi-branch customer, company size and funding.

Who was there

Compassus

Evan Kramer

Leads Innovation & Operational Excellence

Colin Highland

Initiative lead · PT, ~25 years in home health

Michael Thompson

VP, Operational Excellence

Paige Huffman

Senior Project Manager, Operational Excellence

CareStitch

Shelley Ackerman, PT, DPT

COO · led the demo · 10+ years home health PT

Natalie Bell

Customer Operations Manager · 20 years home health, ran centralized HCHB scheduling

Michael

Staff Engineer

Peter Yang

Operations lead · computer scientist and founder

Jason

CTO

Built from the Teams recording (1:59:33) with a machine transcript and extracted screens, plus CareStitch’s questionnaire return where noted. Speaker attribution is inferred from context where the transcript has no labels.

Prepared for Compassus Operational Excellence · Internal · September 2026